

Computer Graphics: Lecture 7

Textures and Bind Groups

Slide Deck © Grant Williams, 2026, License: [CC-BY-SA 4.0](https://creativecommons.org/licenses/by-sa/4.0/)

Welcome Back!

Last time we learned all about color blending and vertex attributes.

We learned that every vertex can have multiple attributes attached to it, and that this data is passed to the vertex shader, one per `@location`.

We learned that we can interleave attributes by putting each attribute in some order, one per vertex.

We learned that the stride of a vertex buffer is how much to add to get from one vertex to another. This is the size of a vertex if attributes are interleaved. If we split attributes into different buffers, this will be the length of one “row”.

This time

We're finally going to learn about textures.

This is an important topic. Textures are the cornerstone of 3D graphics. Not only do they give objects their detail, they can be used as lookup tables for all kinds of effects (shadows, normal maps, reflections, the sky, many more things)

Furthermore, when we render a frame, we're actually drawing into a texture (which is why I've been handwaving `context.getCurrentTexture()`).


What is a texture?

Textures are images that we project on 3D shapes

But this is a bit oversimplified.

In fact, they have a complex internal storage, with multiple levels of images of different sizes.

We will learn the details in the advanced texturing later in the course.



A texture. (by [GGBot](#)—scaled up beyond its intended dimensions)

Why textures?

Textures actually serve a lot of purposes.

Our first use of them will be to give detail to 3D objects.

Since we haven't learned about perspective yet (it's coming, I promise!), textures will be a way to draw pixels on the screen instead of just colored gradients.

They can also be lookup tables, temporary storage, palettes, depth buffers, and even define shadows. They are extremely versatile.

What textures?

Textures can be 1-dimensional, 2-dimensional, or 3-dimensional. We will discuss 2-dimensional textures mainly in this course as those are the most commonly used.¹

In the same way that a normal image is composed of *pixels*, textures are composed of **texels**.

A **texel** is a *TEXTure ELement*.

Why a different term? Because like I mentioned before, textures actually have multiple levels. A texel is a pixel in one of the levels.

¹1D textures are useful for gradients and borders, 3D textures are useful for volumetric fog and medical imaging.

How textures?

In 3D textures are not just images we draw “at an angle”.

Instead, we draw a triangle, but we **sample** the texture as we do so.

Sampling a texture is like looking up a pixel in an image, but it is smooth. We can smoothly look up a point in between several texels.

We do this for every fragment on the triangle.

The key is knowing specifically which texels to look up at which point on the triangle.

Let's get concrete

Textures are a little weird, but once you understand how they work, they gradually become intuitive.

To make them intuitive, we're going to focus on one simple problem: drawing a textured quad.

That is, we're going to draw a rectangle, but we'll modify our fragment shader to look up the texels instead of blending a color.

UV coordinates

Textures are sampled using a special coordinate system. Understanding this coordinate system is critical to getting good at textures.


The coordinate system is called **uv** coordinates.

u and **v** are dimensions, like xy or rgb.

They are a way to describe the position of a sample on a texture, regardless of how big the texture is.

Let's see how they work...

Example UVs




The u, v coordinate system treats the *origin* as the top left.

“ u ” and “ v ” refer to far to go down and right.

Our goal is to choose a point to “sample”, meaning, to retrieve the color at.

If we choose $(0.5, 0.5)$, we are sampling the very center of the image.

Here is a 4-by-4 texture with UV coordinates



Each cell is .25 wide and tall. The center (0.5, 0.5) is between 4 texels. If we sample that point, the GPU could give us one of those 4 colors, or it could average them together depending on our sampling mode.